

The Extremes of Good and Evil

Master Thesis

presented by
Earl Hickey
Matriculation Number 9083894

submitted to the
Data and Web Science Group
Prof. Dr. Right Name Here
University of Mannheim

August 2014

Contents

1	Phase I - Data Translation	1
1.1	Use Case & Data Profiling	1
1.2	Consolidated Schema & Transformations	1
2	Phase II - Identity Resolution	2
2.1	Gold Standard	2
2.2	Matching Rules	3
2.2.1	General Setup	3
2.2.2	Major Challenge - Named Entity Matching	3
2.2.3	Evaluation	4
2.3	Blockers	5
2.4	Analysis of Errors	5
3	Phase III - Data Fusion	6
A	Program Code / Resources	8
B	Phase II - Identity Resolution	9
C	Further Experimental Results	10

List of Figures

List of Tables

2.1	Comparator Overview	3
3.1	Good vs. Evil	6

Chapter 1

Phase I - Data Translation

If you cite something, do it in the following way.

- Conference Proceedings: This problem is typically addressed by approaches for selecting the optimal matcher based on the nature of the matching task and the known characteristics of the different matching systems. Such an approach is described in [?].
- Journal Article: S-Match, described in [?], employs sound and complete reasoning procedures. Nevertheless, the underlying semantic is restricted to propositional logic due to the fact that ontologies are interpreted as tree-like structures.
- Book: According to Euzenat and Shvaiko [?], we define a correspondence as follows.

These are some randomly chosen examples from other works. Take a look at the end of this thesis so see how the bibliography is included.

1.1 Use Case & Data Profiling

1.2 Consolidated Schema & Transformations

Chapter 2

Phase II - Identity Resolution

2.1 Gold Standard

In order to create the gold standard we ran initial identity resolutions with cheap matching rules. With a threshold of 0.3 we used three different matching rules:

- Jaccard-3-Grams on company names (with frequent tokens removed)
- Levensthein Similarity on company names (with frequent tokens removed)
- a combined matching rule of the above (50/50)

The results were then combined into one file which for each individual correspondence outlined the similarity calculated by every similarity measure as well as the company names.

We then labeled the matches into matching record pairs, corner-case matches and non-matches, and non-matches. Correspondences were labeled as certain matches if at least one of the similarity scores had a high matching threshold > 0.95 and an actual match was present. Correspondences were labeled as fuzzy if the similarity measures did not agree on a rating which was indicated by the average similarity, i.e. one similarity measure had a high certainty while the other one had a low certainty. Correspondences with low average similarity were labeled as non-matches after verifying that they indeed did not match. To achieve the distribution according to the rule of thumb outlined in the lecture, which states to include 20% matching record matching pairs, 30% corner-case matches and non-matches (fuzzy), and 50% non-matching record pairs a random sample out of the labeled correspondences was drawn. The data was then split into train and test set using a python script, with a test size of 0.25 and stratified on the gold standard category.

After running several identity resolutions the evaluation logs were analyzed and the gold standard was amended to cater to the identified false positive correspondences (according to the previous gold standard) by adding them to the gold standard, usually as corner-cases.

2.2 Matching Rules

2.2.1 General Setup

The company name was the sole variable we could rely on during the identity resolution. To this end, we implemented several string-based comparators which are outlined in table 2.2.1 and subsequently combined them to form different matching rules. We implemented a plethora of different matching rules, therefore we limit ourselves to the top 10 based on their F1-scores. Every comparator had options to include certain pre- and post-processing steps. Preprocessing steps generally included lowercasing, removal of punctuation, and removal of whitespaces (the latter was omitted for token based similarity metrics). We also implemented post-processing capabilities. These included a threshold after which the similarity was set to zero, and an option to boost or penalize the similarity based on a certain threshold. For example, a similarity might be boosted up using a particular function above a threshold of 0.8 and penalized below.

2.2.2 Major Challenge - Named Entity Matching

As a first challenge in our project we had to find out that company names are inherently difficult to match. We found ourselves facing similar challenges as the

ID	Similarity measure	Parameters	Preprocessing		Focus
			PWL ¹	Rm FT ²	
1	Jaccard on ngrams	n: ngram length	✓	(✓)	Overlap
2	Jaccard on tokens		*	(✓)	Overlap
3	Levensthein		✓	(✓)	Typos / Edit-distance
4	Longest Common Subsequence	Normalization Flag	✓	(✓)	
5	RogueN on Tokens [2]		✓	-	Overlap

Table 2.1: Comparator Overview ¹Lowercasing and removal of punctuation and whitespaces - latter not removed for token-based similarity metrics, ²Removal of frequent tokens, * Preprocessing done by pre-implemented similarity measure, ✓ used in comparator, (✓) optional

Dutch Central Bank ¹. As with named entities in general every data source has a different level of detail, different data quality, and use of abbreviations among others. Regarding this, our matching rules had to cater to the following challenges:

1. **Company Name:** In part the name of the legal entity was used, in other cases the name of the group, and in other cases some abbreviation (e.g., Anheuser-Busch InBev Germany Holding GmbH vs. Anheuser-Busch InBev vs. AB InBev) or tokens of the name were omitted (e.g., Royal Dutch Shell vs. Shell). While all these names refer to the same entity, our matching rules evolved to address this.
2. **Data Quality:** We had to cope with general data quality issues which are represented for example by typos.
3. **Frequent tokens:** There are several tokens that have a higher frequency in company names. These include for example legal entity descriptors (limited, incorporated, ...), industry descriptors (bank, motors, pharmaceuticals, ...) and stop words (the, and, of ...). These may let names seem more similar than they actually are (General Motors vs. Hyundai Motors). We analyzed frequent tokens across our datasets and provided the matching rules with the option to remove frequent tokens.
4. **Token order:** Company names of different companies might be composed of similar tokens in a different order. Token-based similarity metrics alone would classify such names as similar although they are not (Commercial National Financial vs. National Financial Group). This means that token order matters.

We addressed these challenges by combining comparators with different strengths in matching rules which we systematically evaluated at different final matching thresholds. We also implemented some new similarity metrics to cater to our needs (Comparator 4 and 5). We also faced a second challenge concerning dataset sizes which will be outlined in 2.3.

2.2.3 Evaluation

We evaluated local and global matching strategies.

¹<https://medium.com/dnb-data-science-hub/company-name-matching-6a6330710334>

2.3 Blockers

We used three different types of blockers (no blocker, symmetric, sorted neighborhood) and different blocking key generators. The following key generators were implemented:

1. First letter of the company name.
2. Qgrams and first letter of the company name. [1] suggest this blocking technique to ensure a low degree of missed pairs while staying computationally efficient.
3. N starting characters of each company name token.

2.4 Analysis of Errors

Chapter 3

Phase III - Data Fusion

Ontology	Baselines			Decision Tree			
	M(edian)	G(ood)	E(vil)	results	Δ -M	Δ -G	Δ -E
#301	0.825	0.877	0.877	0.855	+0.030	-0.022	-0.022
#302	0.709	0.753	0.753	0.753	+0.044	+0.000	+0.000
#303	0.804	0.860	0.891	0.816	+0.012	-0.044	-0.075
#304	0.940	0.961	0.961	0.967	+0.027	+0.006	+0.006
Average	0.820	0.863	0.871	0.848	+0.028	-0.015	-0.023

Table 3.1: Comparison between the Good and the Evil

Bibliography

- [1] Luis Gravano, Panagiotis G Ipeirotis, H V Jagadish, Nick Koudas, and T Labs. Approximate String Joins in a Database (Almost) for Free. page 10.
- [2] Chin-Yew Lin. ROUGE: A Package for Automatic Evaluation of Summaries. In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain, July 2004. Association for Computational Linguistics.

Appendix A

Program Code / Resources

The source code, a documentation, some usage examples, and additional test results are available at ...

They as well as a PDF version of this thesis is also contained on the CD

Appendix B

Phase II - Identity Resolution

-ROM attached to this thesis.

Appendix C

Further Experimental Results

In the following further experimental results are ...

Ehrenwörtliche Erklärung

Ich versichere, dass ich die beiliegende Master-/Bachelorarbeit ohne Hilfe Dritter und ohne Benutzung anderer als der angegebenen Quellen und Hilfsmittel angefertigt und die den benutzten Quellen wörtlich oder inhaltlich entnommenen Stellen als solche kenntlich gemacht habe. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen. Ich bin mir bewusst, dass eine falsche Erklärung rechtliche Folgen haben wird.

Mannheim, den 31.08.2014

Unterschrift